AI FOR MANAGERS · STUDY GUIDE

The Output Looks Fine — That's the Problem

If the work looks good, what's the problem?

MODULE 2 · LESSON 2 | THE NEW PERFORMANCE PROBLEM | Decision lesson · 6–8 min

THE CORE IDEA

AI errors are dangerous because they often look correct.

Credible-looking output can carry hidden risk — and without a shared way to check AI-assisted work, there is nothing to point to when something goes wrong.

1 Something goes wrong — and you can't trace it

A piece of work goes out from your team — well-structured, clearly written, no obvious problems. Then a client flags a figure that doesn't match their data, or a decision gets made on information that turns out to be incorrect. You go back to understand what happened — and realize you have no idea. You don't know which parts were AI-generated, what was reviewed, whether the sources were checked. There is no trail and no agreed process to point to. Errors slipping past review is a familiar problem; this is its new version. The work looked fine. That was the problem. Underneath sit two problems that together turn a single slip into an ongoing risk.

Problem 1 — the invisible error

AI mistakes that don't look like mistakes.

Problem 2 — the missing net

No shared way to catch them before the work goes out.

2 The invisible error

When humans make mistakes in knowledge work, there are usually surface signals — something is missing, the logic trails off, a number doesn't add up. Your review habits were built to notice those traces. AI introduces a different pattern: work that is structured, fluent, and confident even when something inside it is wrong. The error enters *inside the process*, before your review ever sees the finished work — so the surface arrives polished while a source, figure, assumption, or recommendation underneath it is wrong. A fabricated statistic reads exactly like a real one. A misrepresented source sounds authoritative.

TERM — HALLUCINATION

AI producing output that is factually wrong but presented as if it were correct. The management implication: you cannot use how confident or polished the work sounds to judge whether it is right.

In one study of management consultants, those who used AI on a task built to fall outside its range were about **19 percentage points less likely** than colleagues without AI to reach the right answer — and they finished faster. In practice: an attorney in the US submitted a legal brief citing cases that did not exist, all generated by AI. The citations were plausible and professionally formatted; nothing in the brief signalled that checking them against a legal database was necessary. **The errors AI introduces are often not the ones managers are trained to catch.**

3 The missing net

Even when an error like that slips through, most teams have no net to catch it. AI is being used without a shared understanding of what that means for quality, review, or accountability: hardly anyone has agreed what "reviewed" means when AI contributed, what gets disclosed, or what the standard is for client-facing work. Accountability needs a process – it is not enough for each person to decide privately what counts as checked. **Shadow AI** – the quiet, unannounced use – widens the gap: most people are not hiding anything; no one ever agreed what should be mentioned. The result is a team producing AI-assisted work at volume with no shared way to check it.

When something goes wrong, you can't blame the tool. **Accountability follows whoever sent the work.** The gap is the team's problem – but the person who hit "send" is left holding the consequences. And the gap persists even in compliance-heavy organizations: a policy, an approved tool list, a completed training tell you nothing about whether *this* piece of work was checked the right way. The real question isn't "did we follow the policy?" It's "do we actually know what got checked before this went out?"

4 Where your exposure sits right now

Put the two problems together – errors that don't look like errors, and no net to catch them – and a single slip becomes an ongoing risk. Three questions show where you currently stand – not to solve anything yet, but to see clearly.

THREE QUESTIONS ON YOUR EXPOSURE

1. Do I know which of my team's recent outputs involved AI?
2. If something went wrong in recent work, could I trace what happened and why?
3. Have my team and I ever agreed on what "reviewed" means when AI is part of the process?

If the answer to all three is no, the gap is real and it is open right now. The next lesson builds the picture that lets you answer them.

WHAT TO DO DIFFERENTLY AFTER THIS LESSON

- Stop using polish and confidence as a proxy for correctness – the dangerous errors **read as correct**.
- Assume your review habits were built for human error patterns; AI's errors enter earlier, inside the process.
- Treat "no incident yet" as no evidence – the gap stays invisible until something goes wrong.
- Remember accountability follows whoever **sent** the work, not the tool – a shared checking standard protects your people, not just quality.
- Sit with the three exposure questions honestly – they define the gap this module starts closing.

Module 2 · The New Performance Problem – this module runs the Assess move: replacing assumptions with a baseline of what is actually happening. **Next: Lesson 2.3 – How to Assess Your Team's AI Use** – the method that turns the missing picture into your Team AI Baseline.