AI FOR MANAGERS · STUDY GUIDE

When AI Errors Become Business Decisions

What happens when AI gets it wrong and nobody catches it?

MODULE 4 · LESSON 1 | CONTROL: CONFIDENCE | Decision lesson · 8–10 min

THE CORE IDEA

AI risk becomes real when nobody owns verification.

Caught in the document, an error is an edit. Caught after the decision, it's a commitment you have to live with.

1 From a standard to a control

Align gave you a shared standard for how your team works with AI. But a standard says how work *should* be done — not that it actually is. Closing that gap is the job of **Control**, the framework's third move: Control puts the checks where the risk enters the work. Its output is the **Checkpoint Map** — verification moments placed inside your team's workflows, protecting your data, your decisions, and outside trust. This lesson names the first risk that map has to address.

2 An error that becomes a decision

A piece of work lands in your inbox — a competitive analysis, a briefing, a business case. It reads well; the data points are specific: percentages, figures, named sources. You review it, edit, send it on. Three days later someone chases one of the figures — and it doesn't exist. Neither does the study it was cited from. Caught there, the error is an edit. Now put the discovery three months later, after approval, once resource allocation has begun: the figure justified the investment, commitments are in place, and the error is no longer in a document — it is in a plan. That is where this risk lives: not in the wrong output, but in what happens when it reaches a decision before anyone checks it.

3 Why the error travels

You already know AI produces confidently wrong output with no signal that something is wrong. What matters here is what that output does next: **it moves**. A team member trusts it — nothing in the work gave them a reason not to — and it travels down a chain of hands toward a business case or a leadership presentation, each person assuming the previous one had checked the key claims. Nobody had, because a verification moment was never explicitly placed anywhere in the chain.

That is structural, not bad luck. Fluency and accuracy are independent in AI output, so a wrong answer looks identical to a correct one — and we trust output that looks authoritative: specific, confident, well-formatted. The result is an asymmetry: errors enter the work quickly; their consequences take far longer to unwind.

4 Normal review isn't built for this

The question isn't whether your team is diligent. Normal review catches errors that *look* like errors — poor logic, thin evidence, an argument that doesn't hold together. It is not designed to catch a fabricated statistic that reads exactly like a real one, attributed to a real source. The reframe: **AI-generated information moving toward a decision is a different category of risk** than the same information in a draft or an internal note. Treating them the same is what lets errors travel.

So before work informs a decision, one thing must be clear: who checks the key claims. That doesn't have to be you – your job is to make sure there is a someone, with a name. The Team AI Baseline shows where AI is in the work; the standard's review agreement gives it a second pair of eyes. Neither one names that person.

> **TERM – CHECKPOINT**
>
> A named verification moment inside a workflow, written as **When → What → Who** – the moment, the check, the owner. It is the piece this module adds; the Checkpoint Map is the output of Control.

5 The first diagnostic

Not the Checkpoint Map yet – that comes in Lesson 4.3. For now, one diagnostic move: read any AI-assisted work headed for a decision with three questions in mind.

THREE QUESTIONS TO HOLD IN MIND

1. **What is the key claim this decision rests on?** The one figure, benchmark, or conclusion that, if wrong, would change the decision.
2. **Where did that claim come from – AI, or a verified source?** Open the citation: does the source actually support the claim?
3. **Who verified it before it reached me?** You'll usually have to ask. If the answer is "nobody" – or "I assumed someone had" – that is the gap.

A diagnostic, not a system: start asking these – first of your own incoming work, then of how work moves through the team. Where question three keeps coming back "nobody," you've located a gap the Checkpoint Map will need to close. For now, the move is to notice.

WHAT TO DO DIFFERENTLY AFTER THIS LESSON

- Treat AI-assisted work heading toward a decision as a **different risk category** than drafts and internal notes.
- Assign verification ownership – make sure one **named person** checks the key claims before work informs a decision. You don't verify everything yourself.
- This week, run the three questions on one real document that informed – or could have informed – a decision. Use the actual document and its sources, not your memory of it.
- Write down where the verification moment was – or wasn't. Don't fix anything yet; those notes are raw material for the Checkpoint Map.

Module 4 · Control: Confidence – this module puts the checks where the risk enters the work; its output is the Checkpoint Map.
Next: Lesson 4.2 – AI Risk Goes Beyond Wrong Answers – the risks that aren't about quality at all: your data, your exposure, and the trust of people outside the team.