AI FOR MANAGERS · STUDY GUIDE

AI Risk Goes Beyond Wrong Answers

What else can go wrong beyond bad AI output?

MODULE 4 · LESSON 2 | CONTROL: CONFIDENCE | Decision lesson · 6–8 min

THE CORE IDEA

Output quality is only one category of AI risk.

Risk also enters with what your team puts into AI tools — and with what leaves the team and reaches people outside it.

1 One risk named, two to go

The last lesson named the first risk the Checkpoint Map must address: a wrong answer reaching a decision unchecked stops being a document problem and becomes a commitment you have to live with. This lesson names two more — what goes *into* AI tools, and what comes *out* and reaches people outside the team. Three risks complete the picture of where AI risk actually enters.

2 Risk two — sensitive data goes in before anyone asks

Someone on your team, stuck on a client proposal's executive summary, pastes it into an AI tool for a better opening — it works, they submit on time. They weren't being reckless; they were doing their job. But that one paste may have moved a document carrying your client's name and your proposed approach outside the protections you assume are in place — no warning, and it can repeat quietly, piece after piece.

Most people use AI tools like a search engine: something in, something out, contained. But AI tools are **work systems** — account settings, vendor terms, retention rules — and what happens to entered information depends on a setup most team members never thought to check. And the deeper problem sits before the configuration question: nobody asked whether the document should go in at all. The data agreement in your Team Working Standard already names what goes into AI tools and what doesn't — but **knowing a rule and reaching the moment that triggers it are two different things**. Without that moment, the rule gets applied when people remember it: inconsistently, under pressure, not by default.

3 Risk three — AI errors reach external audiences

The third risk is on the way out. A client finds an error in a deliverable — a figure that doesn't match the source data, traced to an unchecked AI-generated section. You correct it, they accept. Then, at the next kickoff, a new request: can drafts come two days earlier, so they can run their own fact-check? That is the reputational consequence. Not a lost contract — a quiet shift: the client now plans around the chance that your work is wrong, and you'll probably never hear about it again.

This risk is different in kind. An error that stays inside the team is **operational** — correctable, revisitable. An error that reaches people outside it is **relational** — trust is slower to rebuild than a document is to fix. And it's broader than accuracy: external work can also fail by not fitting the client, the relationship, or the context. The review agreement in your standard already names which work needs human judgment before it's used or shared; what no rule names is the moment — where that review happens, by whom, checking what.

4 Three risks, three controls

Three risks, entering at different moments, causing different kinds of damage – and the rules for each already exist. But rules tell people *what* to do, not *when*. And *when* is where the risk enters.

Decisions
AI errors heading for a decision need a **verification moment** before the work gets there.

Data
Sensitive data entering AI tools needs a **check before content goes in** – not after.

External output
Work going to clients, partners, other departments, or public audiences needs a **different review standard** than internal work – because the consequence of error is different.

These aren't AI's only risks – capability and fairness are people questions the course returns to later; these three are the ones a checkpoint can directly guard against. And a checkpoint is small: a named moment, a minutes-long check, one owner – one question before pasting, one source opened before approval, one sharper read before it leaves. That is the shift from a rule the team knows to a control that holds.

5 The checkpoint scan

This week, run this on the places you already know AI sits in the team's work – the Baseline, the mapped workflows.

SCAN THREE AREAS – IS THERE A CHECK, AND IS IT NAMED?

1. **Decisions.** Where AI-assisted work feeds a decision: is anyone named to verify the key claims before it gets there?
2. **Data.** Where sensitive material could go into a tool: is there a moment where someone asks "should this go in?" – or does it rely on people remembering?
3. **External output.** Where work leaves the team: does it get a sharper check than internal work – or the same one?

Nothing gets fixed yet. The task is finding the spots where the answer is "nobody" or "nothing" – those gaps are the raw material the next lesson turns into the Checkpoint Map.

WHAT TO DO DIFFERENTLY AFTER THIS LESSON

- Stop treating output quality as the whole risk picture – watch **what goes into tools** and **what leaves the team**, too.
- Hold external work to a sharper review than internal work – internal errors are operational, external ones are relational.
- When you spot a risk, resist the instinct to just re-tighten the rules – place a **moment** where the rule takes effect.
- Run the three-area scan this week and note every check that has no name on it.

Module 4 · Control: Confidence – this module puts the checks where the risk enters the work; its output is the Checkpoint Map.
Next: Lesson 4.3 – Build and Activate the Checkpoint Map – turning the gaps you've found into named checkpoints that hold inside real work.